

ECE30030/ITP30010 Database Systems

SQL DDL

Reading: Chapter 3

Charmgil Hong

charmgil@handong.edu

Spring, 2023

Handong Global University

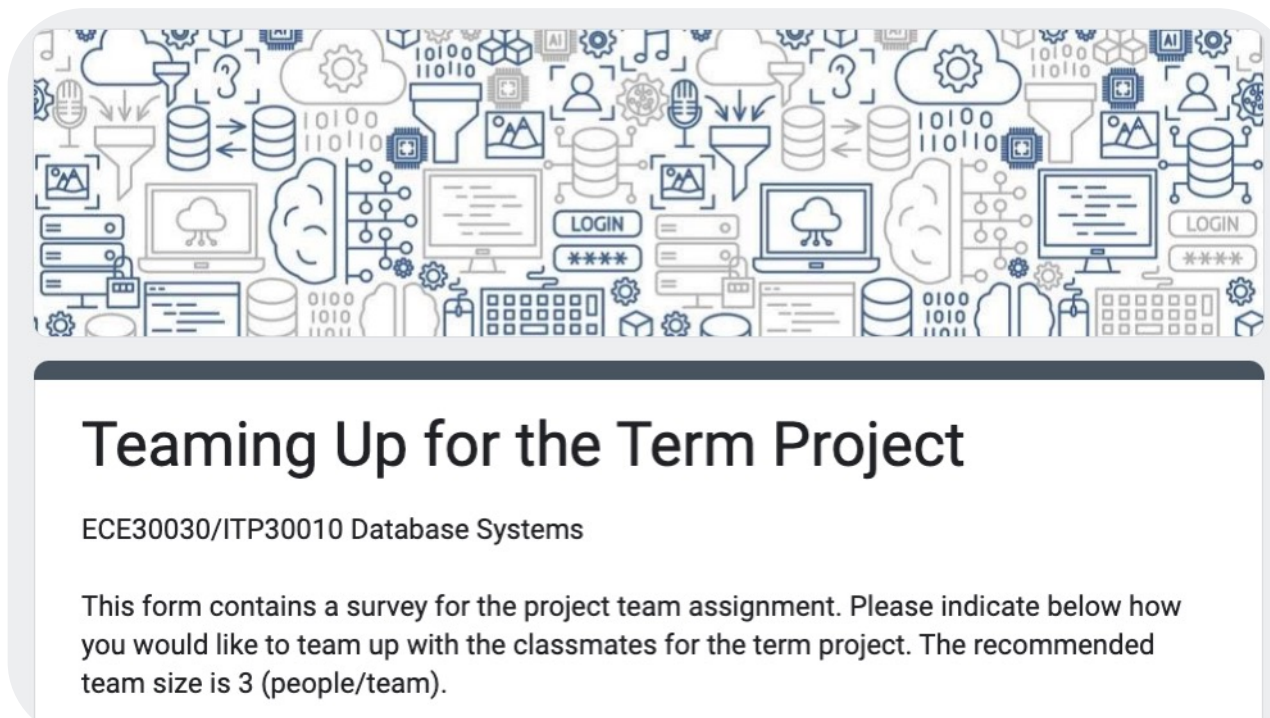


Announcement

- Homework assignment #3 is out
 - Due: By the end of Saturday, April 22
 - Please start early

Announcement

- Forming teams for the term project
 - Response due: Monday, April 10
 - URL: <https://forms.gle/kQWG9ML6fqytYm7p7>
 - Problem & data release: Week #8

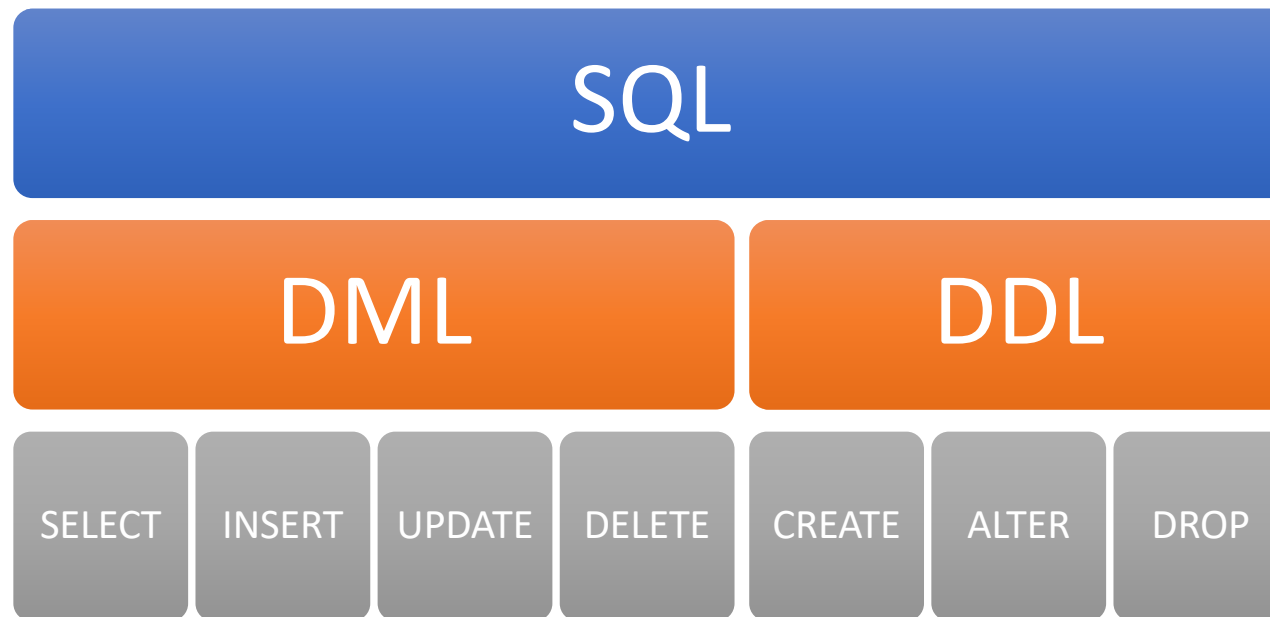


Teaming Up for the Term Project

ECE30030/ITP30010 Database Systems

This form contains a survey for the project team assignment. Please indicate below how you would like to team up with the classmates for the term project. The recommended team size is 3 (people/team).

SQL Commands



Data Definition Language

- The SQL data-definition language (DDL) allows the specification of information about relations, including:
 - The **schema** for each relation
 - The **type** of values associated with each attribute
 - The Integrity **constraints**
 - The set of **indices** to be maintained for each relation
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk
- Three key commands
 - CREATE
 - ALTER
 - DROP

CREATE DATABASE

- To initialize a new database

- Basic syntax:

CREATE DATABASE *database_name*


- One can specify the default character encoding method along with this command

- **CREATE DATABASE** test
DEFAULT CHARACTER SET utf8 *utf8 mb4*
COLLATE utf8_unicode_ci;

- Collation: a set of rules that defines how to compare and sort character strings

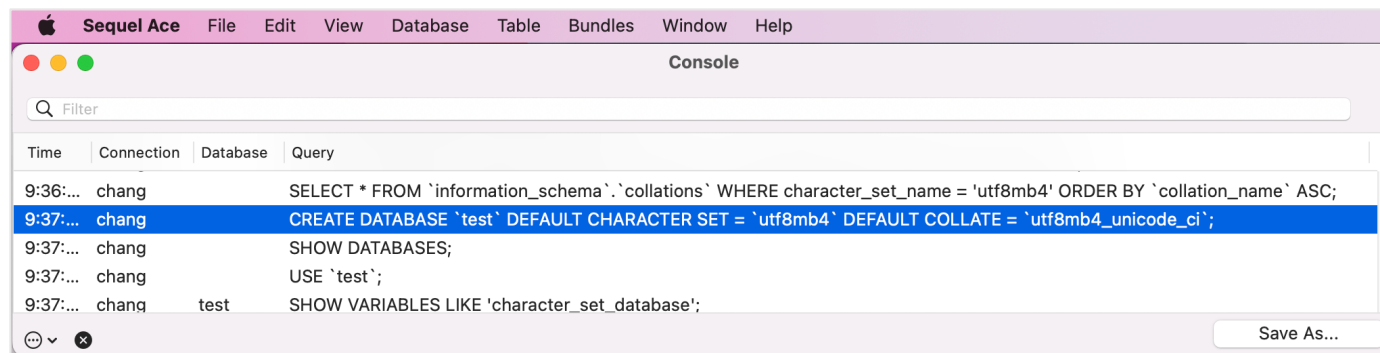
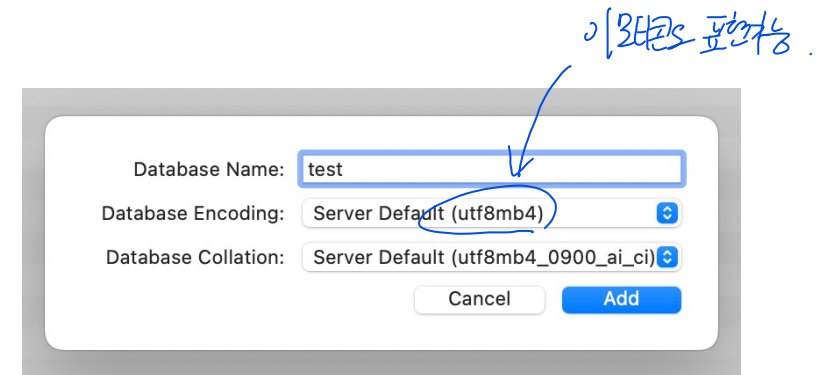
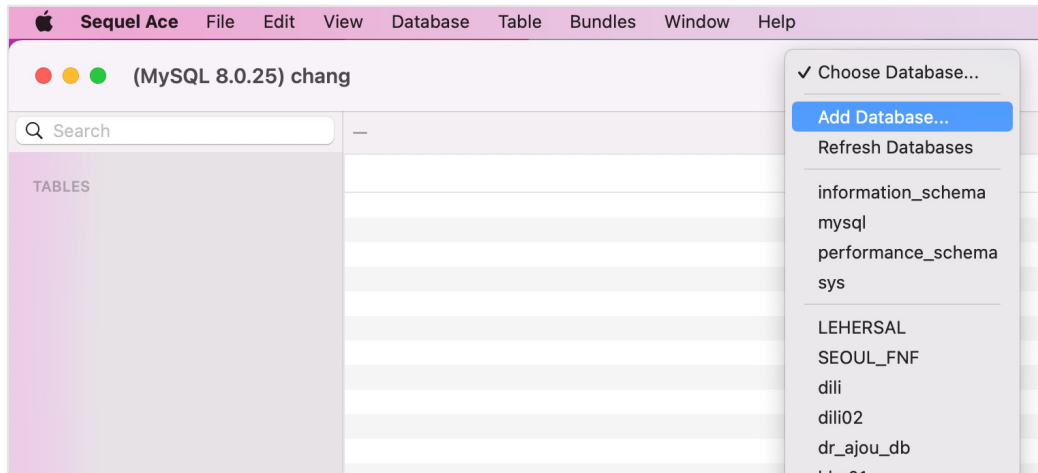
- After creating a database, to use it
USE *database_name*

- Ref: <https://dev.mysql.com/doc/refman/8.0/en/charset-charsets.html>



Example: Creating a Database on Sequel Pro

- Creating a new database



Agenda

- SQL DDL (Data Definition Language)

CREATE TABLE

- To create a new table

- Basic syntax:

```
CREATE TABLE table_name(  
    Col1_name      data_type[(size)],  
    Col2_name      data_type[(size)]  
)
```

- E.g., Creating a table with four columns

- **CREATE TABLE** *books*(
 ISBN CHAR(20),
 Title CHAR(50),
 AuthorID INTEGER,
 Price FLOAT)

Data Types in SQL

- Following categories of data types exist in most DBMSs
 - String data
 - Numeric data
 - Temporal data
 - Large objects

String Data in SQL

- SQL Data Types

- **CHAR(*n*)**: Fixed length character string, with user-specified length *n*
 - Maximum length *n* = [0, 255]
- **VARCHAR(*n*)**: Variable length character strings, with user-specified maximum length *n* *it just specifies the maximum length. => n may not be the same as assigned storage.*
 - Maximum length *n* = [0, 65,535]
- *If the length is always the same, use a CHAR-type attribute;
if you are storing wildly variable length strings, use a VARCHAR-type attribute*

off records:

- **TEXT**: for strings longer than the range of VARCHAR

- TINYTEXT 0 – 255 bytes
- TEXT 0 – 65,535 bytes
- MEDIUMTEXT 0 – 16,777,215 bytes
- LONGTEXT 0 – 4,294,967,295 bytes

*these data type is not stored in a table.
but in the external storage.
So it links to the actual space in the external storage.
which means it takes more time to retrieve*

String Data in SQL

- Difference between CHAR and VARCHAR

Value	CHAR(4)	Storage	VARCHAR(4)	Storage
''	' '	4 bytes	''	1 bytes
'ab'	'ab '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefg'	'abcd'	4 bytes	'abcd'	5 bytes

Numeric Data in SQL

- SQL Data Types
 - **INT, INTEGER:** *Integer* (a finite subset of the integers that is machine-dependent)
 - **SMALLINT:** *Short integer* (a machine-dependent subset of the integer domain type)
 - **BIGINT:** *Long integer* (a machine-dependent subset of the integer domain type)
- **TINYINT** and **MEDIUMINT** are also available

Numeric Data in SQL

- Different R-DBMSs support different combinations of those integer types

	Bytes	MySQL	MS SQL	PostgreSQL	DB2
TINYINT	1	✓	✓		
SMALLINT	2	✓	✓	✓	✓
MEDIUMINT	3	✓			
INT/INTEGER	4	✓	✓	✓	✓
BIGINT	8	✓	✓	✓	✓

- *C.f.*, Oracle only has a NUMBER datatype

Numeric Data in SQL

- SQL Data Types
 - **NUMERIC(p,d)**: Fixed point number (exact value) with user-specified precision of p digits, with d digits to the right of decimal point
 - *E.g.*, **NUMERIC**(3,1) allows 44.5 to be stored exactly, but not 444.5 or 0.32)
 - In MySQL, **DECIMAL** is NUMERIC
 - **FLOAT**: Floating point number (approximate) with **single-precision**
 - **REAL, DOUBLE**: Floating point number (approximate) with double-precision

Numeric Data in SQL

- DECIMAL vs INT/FLOAT/DOUBLE
 - FLOAT and DOUBLE are faster than DECIMAL
 - DECIMAL values are exact
 - Example

floats: FLOAT	decimals: DECIMAL(3,2)
1.1	1.10
1.1	1.10
1.1	1.10

- SELECT SUM(...) → **DECIMAL values are precise**

SUM(floats)	SUM(decimals)
3.3000000715255737	3.30

even though you want 3.30 but
in floating digits format
it actually represent that number.

Temporal Data in SQL

- SQL Data Types
 - **DATE**: 'YYYY-MM-DD'
 - Range: 1000-01-01 to 9999-12-31
 - *E.g.*, '2020-03-01' for March 1, 2020
 - **TIME**: 'HH:MM:SS'
 - Range: -838:59:59 to 838:59:59
 - *E.g.*, '14:30:03.5' for 3.5 seconds after 2:30pm
 - **DATETIME**: 'YYYY-MM-DD HH:MM:SS'
 - Range: 1000-01-01 00:00:00 to 9999-12-31 23:59:59
 - **YEAR**: 'YYYY'
 - Range: 1901 to 2155, or 0000 (illegal year values are converted to 0000)

Temporal Data in SQL

- SQL Data Types
 - **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)
 - A way to track time as a running total of seconds
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - URL: <https://time.is/Unix>

**UNIX
TIME
SINCE 1970**

Temporal Data in SQL

- SQL Data Types
 - **TIMESTAMP(n)**: **Unix time** (time since Jan 1, 1970)
 - A way to track time as a running total of seconds
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - URL: <https://time.is/Unix>

Binary : 01111111 11111111 11111111 11110000

Decimal : 2147483632

Date : 2038-01-19 03:13:52 (UTC)

Date : 2038-01-19 03:13:52 (UTC)



Temporal Data in SQL

- SQL Data Types
 - **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)
 - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
 - Typically used for logging (keeping records of all the system events)
 - Depending on size *n*, the **display pattern** changes

	Format
TIMESTAMP(14)	YYYYMMDDHHMMSS
TIMESTAMP(12)	YYMMDDHHMMSS
TIMESTAMP(10)	YYMMDDHHMM
TIMESTAMP(8)	YYYYMMDD
TIMESTAMP(6)	YYMMDD
TIMESTAMP(4)	YYMM
TIMESTAMP(2)	YY

Large Objects in SQL

- SQL Data Types

- ✓ **BINARY(*n*)**: binary byte data type, with user-specified length *n*
 - Contains a byte strings (rather than a character string)
 - Maximum length *n* = [0, 255]
- ✓ **VARBINARY(*n*)**: binary byte data type, with user-specified maximum length *n*
 - Maximum length *n* = [0, 65,535]
- **BLOB**: Binary Large Object data type
 - TINYBLOB 0 – 255 bytes
 - BLOB 0 – 65,535 bytes (65 KB)
 - MEDIUMBLOB 0 – 16,777,215 bytes (16 MB)
 - LONGBLOB 0 – 4,294,967,295 bytes (4 GB)

On-record

Off-record

CREATE TABLE Construct

- A new relation is defined using the **CREATE TABLE** command:

```
CREATE TABLE r  
    (A1 D1, A2 D2, ..., An Dn,  
    (integrity-constraint1),  
    ...,  
    (integrity-constraintk))
```

- *r* is the name of the relation
 - Each *A_i* is an attribute name in the schema of relation *r*
 - Each *D_i* is the data type of values in the domain of attribute *A_i*
-
- Example: **CREATE TABLE** instructor(
 ID CHAR(5),
 name VARCHAR(20),
 dept_name VARCHAR(20),
 salary NUMERIC(8,2))

Integrity Constraints in CREATE TABLE

- SQL prevents any update to the database that violates an **integrity constraint**
 - Integrity constraints allow us to specify what data makes sense for us
- Types of integrity constraints
 - Primary key: **PRIMARY KEY** (A_1, \dots, A_n)
 - Foreign key: **FOREIGN KEY** (A_m, \dots, A_n) **REFERENCES** r
 - Unique key: **UNIQUE** (A_1, \dots, A_n)
 - Not null: **NOT NULL**
 - Value constraints: **CHECK** (*constraint*), **DEFAULT**

- Example:

```
CREATE TABLE instructor(  
    ID CHAR(5),  
    name VARCHAR(20) NOT NULL,  
    dept_name VARCHAR(20),  
    salary NUMERIC(8, 2),  
    PRIMARY KEY (ID),  
    FOREIGN KEY (dept_name) REFERENCES department);
```

Declaring KEY and UNIQUE Constraints

- An attribute or list of attributes may be declared as PRIMARY KEY or UNIQUE
 - Meaning: no two tuples of the relation may agree in all the attribute(s) on the list
 - That is, the attribute(s) do(es) **not allow duplicates** in values
 - PRIMARY KEY/UNIQUE can be used as an **identifier for each row**
 - Comparison: PRIMARY KEY vs UNIQUE

PRIMARY KEY	UNIQUE
Used to serve as a unique identifier for each row in a relation	Uniquely determines a row which is not primary key
Cannot accept NULL	Can accept NULL values (some DBMSs accept only one NULL value)
A relation can have only one primary key	A relation can have more than one unique attributes
Clustered index	Non-clustered index

Examples

- **CREATE TABLE** *student* (
 ID **VARCHAR(5),**
 name **VARCHAR(20) NOT NULL,**
 dept_name **VARCHAR(20),**
 tot_cred **NUMERIC(3,0),**
 PRIMARY KEY (*ID*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

Examples

- **CREATE TABLE** *student* (
 ID **VARCHAR(5) PRIMARY KEY**,
 name **VARCHAR(20) NOT NULL**,
 dept_name **VARCHAR(20)**,
 tot_cred **NUMERIC(3,0)**,
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *takes* (
 ID **VARCHAR**(5),
 course_id **VARCHAR**(8),
 sec_id **VARCHAR**(8),
 semester **VARCHAR**(6),
 year **NUMERIC**(4,0),
 grade **VARCHAR**(2),
 PRIMARY KEY (*ID, course_id, sec_id, semester, year*),
 FOREIGN KEY (*ID*) **REFERENCES** *student*,
 FOREIGN KEY (*course_id, sec_id, semester, year*)
 REFERENCES *section*);

More Examples

- **CREATE TABLE** *course* (
 course_id **VARCHAR**(8),
 title **VARCHAR**(50),
 dept_name **VARCHAR**(20),
 credits **NUMERIC**(2,0),
 PRIMARY KEY (*course_id*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *course* (
 course_id **VARCHAR**(8),
 title **VARCHAR**(50),
 dept_name **VARCHAR**(20) **DEFAULT** 'Comp. Sci',
 credits **NUMERIC**(2,0),
 PRIMARY KEY (*course_id*),
 FOREIGN KEY (*dept_name*) **REFERENCES** *department*);

More Examples

- **CREATE TABLE** *neighbors*(
 name **CHAR(30) PRIMARY KEY**,
 addr **CHAR(50) DEFAULT '123 Sesame St.'**,
 phone **CHAR(16));**
- Inserting Elmo is a neighbor (inserted with the default value):
 - **INSERT INTO** *neighbors* (*name*)
 VALUES ('Elmo');

name	addr	phone
'Elmo'	'123 Sesame St.'	NULL

More Examples

- **CREATE TABLE** *neighbors*(
 name **CHAR(30) PRIMARY KEY**,
 addr **CHAR(50) DEFAULT '123 Sesame St.'**,
 phone **CHAR(16) NOT NULL**);
- Inserting Elmo is a neighbor:
 - **INSERT INTO** *neighbors* (*name*)
 VALUES ('Elmo');
 - ➔ If *phone* were NOT NULL, this insertion would have been **rejected**

Integrity Constraints Recap

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
 - A single or multiple columns can be specified as a key
 - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0),  
    name             VARCHAR(20),  
    city             VARCHAR(20),  
    state            CHAR(2),  
);
```


Integrity Constraints Recap

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
 - A single or multiple columns can be specified as a key
 - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0),  
    name             VARCHAR(20),  
    city             VARCHAR(20),  
    state            CHAR(2),  
    UNIQUE(name)  
);
```

Integrity Constraints Recap

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
 - A single or multiple columns can be specified as a key
 - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0),  
    name             VARCHAR(20),  
    city             VARCHAR(20),  
    state            CHAR(2),  
    UNIQUE(name),  
    UNIQUE(city, state),  
);
```

Integrity Constraints Recap

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
 - A single or multiple columns can be specified as a key
 - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0),  
    name              VARCHAR(20),  
    city              VARCHAR(20),  
    state             CHAR(2),  
    PRIMARY KEY(ID),  
    UNIQUE(name),  
    UNIQUE(city, state),  
);
```

Integrity Constraints Recap

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
 - A single or multiple columns can be specified as a key
 - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0) PRIMARY KEY,  
    name             VARCHAR(20) UNIQUE,  
    city             VARCHAR(20),  
    state            CHAR(2),  
    UNIQUE(city, state),  
);
```

Integrity Constraints Recap

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
 - A single or multiple columns can be specified as a key
 - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0) PRIMARY KEY,  
    name             VARCHAR(20) UNIQUE,  
    city             VARCHAR(20),  
    state            CHAR(2),  
    UNIQUE(city, state),  
    FOREIGN KEY (state) REFERENCES states  
);
```

Integrity Constraints Recap

- **NOT NULL** – disallowing null values
 - Null values indicate that the data is not known
 - These can cause problems in querying database
 - The Primary Key columns automatically prevent null being entered
 - *C.f.*, **NULL** – can be used to explicitly allow null values

```
CREATE TABLE studio (  
    ID                NUMERIC(5,0) PRIMARY KEY,  
    name             VARCHAR(20) NOT NULL,  
    city             VARCHAR(20) NULL,  
    state            CHAR(2) NOT NULL  
);
```

Integrity Constraints Recap

- **DEFAULT** – A default value can be inserted in any column with this keyword

- *E.g.*, **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) NOT NULL,**
 release_date **DATE DEFAULT sysdate NULL,**
 genre **VARCHAR(20) DEFAULT 'Comedy'**
 CHECK genre IN ('Comedy', 'Action', 'Drama')
)

- In MySQL,
 - **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) NOT NULL,**
 release_date **DATE DEFAULT CURRENT_TIMESTAMP NULL,**
 genre **VARCHAR(20) DEFAULT 'Comedy'**
 CHECK genre IN ('Comedy', 'Action', 'Drama')
)

Integrity Constraints Recap

- **CHECK** – Allows the inserted value to be checked
 - *E.g.*, **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) PRIMARY KEY,**
 release_date **DATE,**
 budget **INTEGER CHECK (budget > 50000)**
)
 - Table-level constraints can be defined; *E.g.*,
 - **CREATE TABLE** *movies*(
 movie_title **VARCHAR(40) PRIMARY KEY,**
 release_date **DATE,**
 budget **INTEGER CHECK (budget > 50000),**
 CONSTRAINT *release_date_const*
 CHECK (release_date BETWEEN '01-Jan-2000' AND '31-Dec-2009')
)

Table Updates (Updating Table Schemas)

- **DROP:** Used to remove elements from a database, such as tables
 - **DROP TABLE r**
 - Remove relation r
 - C.f., **TRUNCATE (TABLE) r** is used to delete the data inside a table, but not the table itself
- **ALTER:** Used to make changes to the table schema
 - **ALTER TABLE r ADD A D**
 - A is the name of the **new attribute** to add to relation r ; D is the **domain** of A
 - All **existing tuples in the relation are assigned *null*** as the value for the new attribute
 - **ALTER TABLE r DROP A**
 - A is the name of an **attribute** in r
 - Dropping of attributes not supported by many databases (MySQL does)

Table Updates (Updating Table Schemas)

- Examples
 - **DROP TABLE** time_slot_backup;
 - **ALTER TABLE** time_slot_backup **ADD** remark VARCHAR(20);
 - **ALTER TABLE** time_slot_backup **MODIFY** remark CHAR(20);
 - **ALTER TABLE** time_slot_backup **DROP** remark;



Table Updates (Updating Table Schemas)

- Examples
 - Drop a column
 - **ALTER TABLE salary**
DROP COLUMN instructor;
 - DROP a PRIMARY KEY Constraint
 - **ALTER TABLE instructor**
DROP PRIMARY KEY;
 - DROP a FOREIGN KEY Constraint
 - **ALTER TABLE instructor**
DROP FOREIGN KEY instructor_ibfk_1;
 - DROP DEFAULT
 - **ALTER TABLE student**
ALTER tot_cred DROP DEFAULT;



Table Updates (Updating Table Schemas)

- Examples
 - Drop a database
 - **DROP DATABASE** university;




Table Updates (Updating Tuples)

- **INSERT**
 - **INSERT INTO** *instructor*
VALUES ('10211', 'Smith', 'Biology', 66000)
- **DELETE**
 - **DELETE FROM** *student*
 - Remove all tuples from the *student* relation
 - **TRUNCATE TABLE** *student*



SQL Commands



INSERT

- Add a new tuple to *course*
 - **INSERT INTO** *course*
VALUES ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
- or equivalently
 - **INSERT INTO** *course* (*course_id*, *title*, *dept_name*, *credits*)
VALUES ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
- Add a new tuple to *student* with *tot_creds* set to null
 - **INSERT INTO** *student*
VALUES ('3003', 'Green', 'Finance', *null*);

INSERT

- Inserting results of other SELECT query
 - Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000
 - **INSERT INTO** *instructor*
 SELECT *ID, name, dept_name, 18000*
 FROM *student*
 WHERE *dept_name = 'Music' AND total_cred > 144;*
 - The **SELECT FROM WHERE** statement is evaluated fully **before** any of its results are inserted into the relation
 - Otherwise queries like
 INSERT INTO table1 SELECT * FROM table1
 would cause problem

UPDATE

- Basic syntax
 - Updating a table
 - **UPDATE** *tablename*
SET *col1_name = new_col1_value, col2_name = new_col2_value, ...;*
 - Updating a table with conditions
 - **UPDATE** *tablename*
SET *col1_name = new_col1_value, col2_name = new_col2_value, ...*
WHERE *predicate;*

UPDATE

- Give a 5% salary raise to all instructors
 - **UPDATE** *instructor*
SET *salary* = *salary* * 1.05
- Give a 5% salary raise to those instructors who earn less than 70000
 - **UPDATE** *instructor*
SET *salary* = *salary* * 1.05
WHERE *salary* < 70000;
- Give a 5% salary raise to instructors whose salary is less than average
 - **UPDATE** *instructor*
SET *salary* = *salary* * 1.05
WHERE *salary* < (**SELECT** **AVG**(*salary*) **FROM** *instructor*);

UPDATE

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
 - Write two UPDATE statements:
UPDATE *instructor*
SET *salary* = *salary* * 1.03
WHERE *salary* > 100000;

UPDATE *instructor*
SET *salary* = *salary* * 1.05
WHERE *salary* <= 100000;
- The order is important
- Can be done better using the **case** statement (next slide)

CASE Statement for Conditional Update

- The following query is equivalent to the previous UPDATE queries
 - **UPDATE** *instructor*
 SET *salary* = **CASE**
 WHEN *salary* <= 100000 **THEN** *salary* * 1.05
 ELSE *salary* * 1.03
 END

UPDATE with Scalar Subqueries

- Recompute and update *tot_creds* value for all students
 - **UPDATE** *student S*
SET *tot_cred* = (**SELECT SUM**(*credits*)
FROM *takes, course*
WHERE *takes.course_id* = *course.course_id* **AND**
S.ID = *takes.ID* **AND**
takes.grade <> 'F' **AND**
takes.grade **IS NOT NULL**);

DELETE

- Basic syntax
 - To remove specific rows
 - **DELETE FROM** *tablename*
WHERE *predicate*;
 - To remove all rows
 - **DELETE FROM** *tablename*;
 - This is equivalent to **TRUNCATE**:
TRUNCATE (TABLE) *tablename*;
 - One cannot truncate a table with foreign key constraints
 - Must disable the constraints first (we will cover **ALTER** when we study SQL DDL):
ALTER TABLE *tablename*
DISABLE CONSTRAINT *constraint_name*;

DELETE

- Delete all instructors
 - **DELETE FROM** *instructor*;
- Delete all instructors from the Finance department
 - **DELETE FROM** *instructor*
WHERE *dept_name*= 'Finance';
- Delete all tuples in the instructor relation for those instructors associated with a department located in the Watson building
 - **DELETE FROM** *instructor*
WHERE *dept name* **IN** (**SELECT** *dept name*
FROM *department*
WHERE *building* = 'Watson');

DELETE

- Delete all instructors whose salary is less than the average salary of instructors
 - Example: **DELETE FROM** *instructor*
WHERE *salary* < (**SELECT AVG** (*salary*)
FROM *instructor*);
- Issue: as we delete tuples from *instructor*, the average salary changes
 - Solution used in SQL:
 1. First, compute **AVG**(*salary*) and find all tuples to delete
 2. Next, delete all tuples found above (**without recomputing AVG** or retesting the tuples)

EOF

- Coming next:
 - Designing a database