

ECE30030/ITP30010 Database Systems

# SQL DDL

*Reading: Chapter 3*

---

***Charmgil Hong***

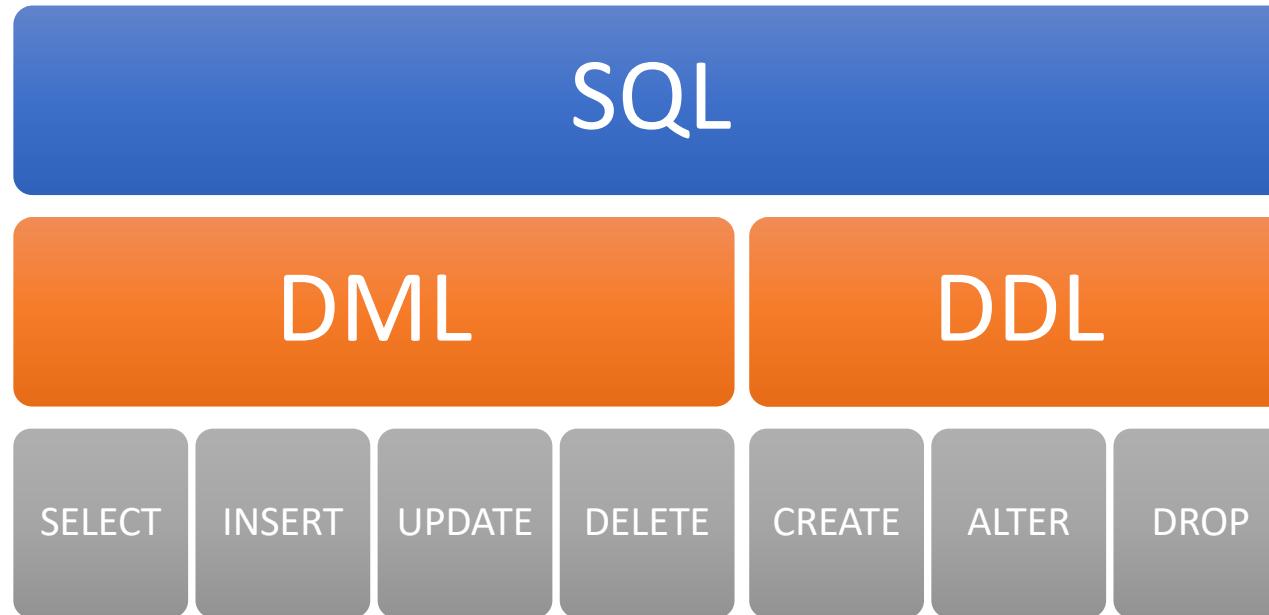
charmgil@handong.edu

Spring, 2023  
Handong Global University



# SQL Commands

---



# Data Definition Language

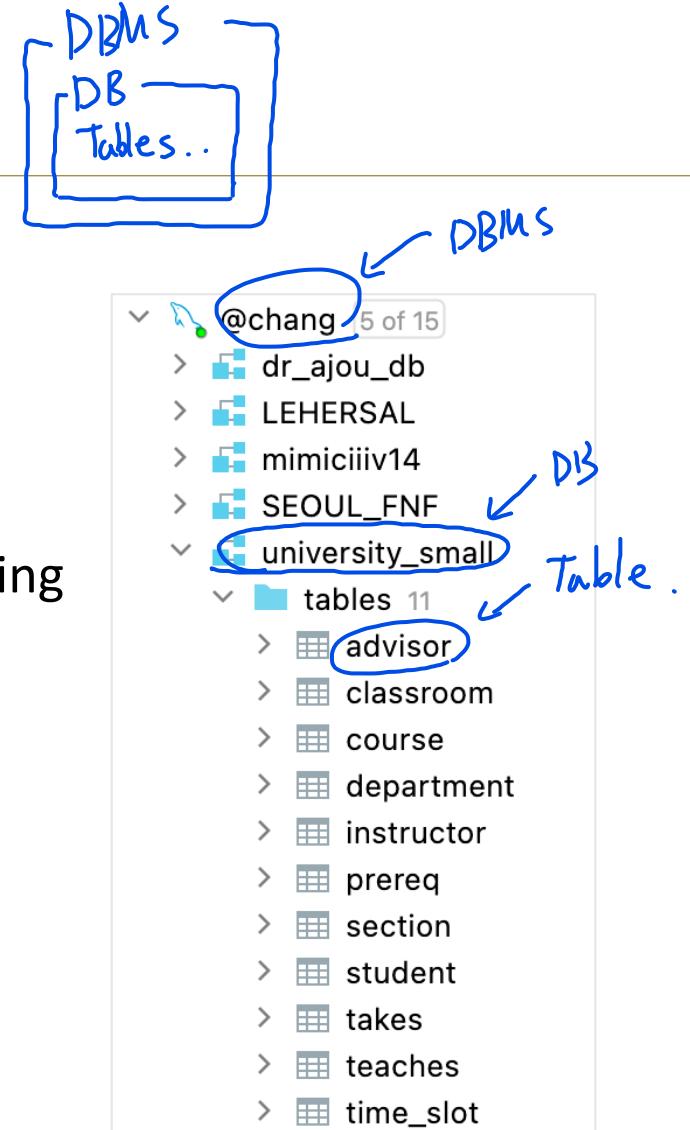
- The SQL data-definition language (DDL) allows the specification of information about relations, including:
    - The **schema** for each relation *Table structure*.
    - The **type** of values associated with each attribute
    - The Integrity **constraints**
    - The set of **indices** to be maintained for each relation
    - Security and authorization information for each relation
    - The physical storage structure of each relation on disk
  - Three key commands
    - CREATE
    - ALTER
    - DROP
- Things you can do with DDL.*
- an additional data structure along with the table (associate) to increasing a performance in reading data.*

# CREATE DATABASE

- To initialize a new database
- Basic syntax:

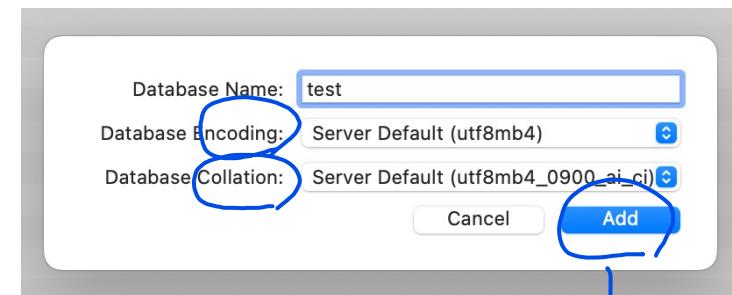
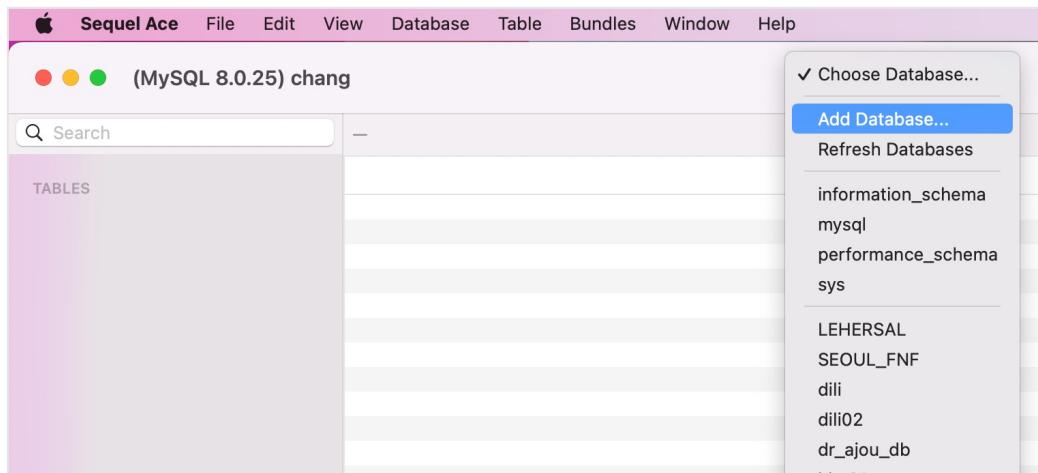
**CREATE DATABASE** *database\_name*.

- One can specify the default character encoding method along with this command
  - **CREATE DATABASE** test  
**DEFAULT CHARACTER SET** utf8  
**COLLATE** utf8\_unicode\_ci;
    - Collation: a set of rules that defines how to compare and sort character strings
  - After creating a database, to use it  
**USE** *database\_name*
- Ref: <https://dev.mysql.com/doc/refman/8.0/en/charset-charset.html>

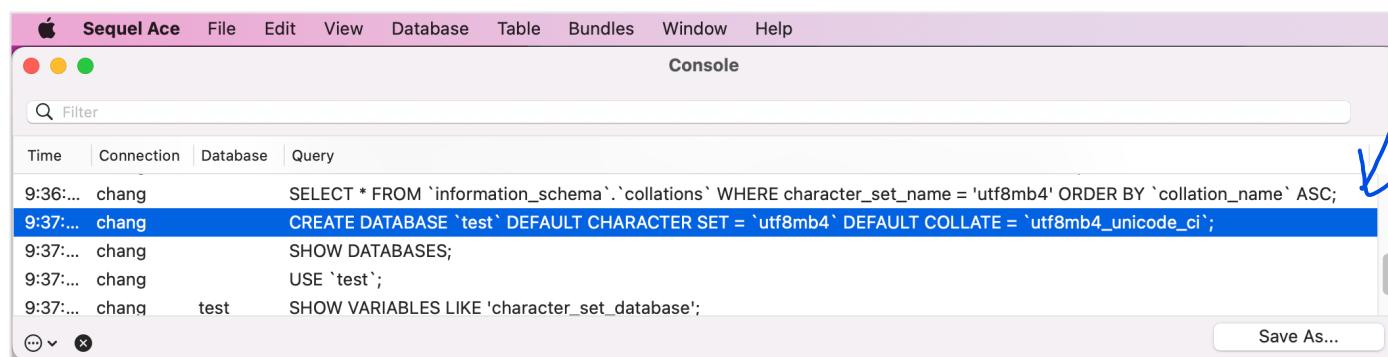


# Example: Creating a Database on Sequel Pro

- Creating a new database



Automatically  
make query  
for it.



```
9:36:...  chang      SELECT * FROM `information_schema`.`collations` WHERE character_set_name = 'utf8mb4' ORDER BY `collation_name` ASC;
9:37:...  chang      CREATE DATABASE `test` DEFAULT CHARACTER SET = `utf8mb4` DEFAULT COLLATE = `utf8mb4_unicode_ci`;
9:37:...  chang      SHOW DATABASES;
9:37:...  chang      USE `test`;
9:37:...  chang      test      SHOW VARIABLES LIKE 'character_set_database';

```

# Agenda

---

- SQL DDL (Data Definition Language)

# CREATE TABLE

---

- To create a new table

- Basic syntax:

```
CREATE TABLE table_name(  
    Col1_name      data_type[(size)],  
    Col2_name      data_type[(size)]  
)      name          datatype .
```

- E.g., Creating a table with four columns

- **CREATE TABLE** books()

```
    ISBN           CHAR(20),  
    Title          CHAR(50),  
    AuthorID       INTEGER,  
    Price          FLOAT)
```

# Data Types in SQL

---

- Following categories of data types exist in most DBMSs
  - String data
  - Numeric data
  - Temporal data
  - Large objects,

# String Data in SQL

## • SQL Data Types

- **CHAR( $n$ )**: **Fixed length character string**, with user-specified length  $n$ 
    - Maximum length  $n = [0, 255]$
  - **VARCHAR( $n$ )**: **Variable length character strings**, with user-specified maximum length  $n$ 
    - Maximum length  $n = [0, 65,535]$
- allocate a fixed size of space.  
regardless of data type.*
- allocate the storage according to the size of data value..*
- If the length is always the same, use a CHAR-type attribute;  
if you are storing wildly variable length strings, use a VARCHAR-type attribute.*

## • **TEXT** for strings longer than the range of VARCHAR.

- TINYTEXT 0 – 255 bytes
- TEXT 0 – 65,535 bytes
- MEDIUMTEXT 0 – 16,777,215 bytes
- LONGTEXT 0 – 4,294,967,295 bytes

16MB.  
4GB.

→ stored in storage  
not in a tree.

# String Data in SQL

- Difference between CHAR and VARCHAR

Value	CHAR(4)	Storage	VARCHAR(4)	Storage
''	' ''	4 bytes	''	1 bytes
'ab'	'ab '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefg'	'abcd'	4 bytes	'abcd'	5 bytes

for null char.  
to indicate  
the end of the  
string.

# Numeric Data in SQL

---

- SQL Data Types
  - **INT, INTEGER**: **Integer** (a finite subset of the integers that is machine-dependent) 32-bit
  - **SMALLINT**: **Short integer** (a machine-dependent subset of the integer domain type) 16-bit
  - **BIGINT**: **Long integer** (a machine-dependent subset of the integer domain type) 64-bit
  - **TINYINT** and **MEDIUMINT** are also available

# Numeric Data in SQL

- Different R-DBMSs support different combinations of those integer types

	Bytes	MySQL	MS SQL	PostgresSQL	DB2
TINYINT	1	✓	✓		
SMALLINT	2	✓	✓	✓	✓
MEDIUMINT	3	✓			
INT/INTEGER	4	✓	✓	✓	✓
BIGINT	8	✓	✓	✓	✓

- C.f.*, Oracle only has a NUMBER datatype

# Numeric Data in SQL

- SQL Data Types *or Decimal*.

- **NUMERIC**(*p,d*): Fixed point number (exact value) with user-specified precision of *p* digits, with *d* digits to the right of decimal point

- E.g., **NUMERIC**(3,1) allows 44.5 to be stored exactly, but not 444.5 or 0.32)

- In MySQL, **DECIMAL** is **NUMERIC**

- **FLOAT**: Floating point number (approximate) with single-precision

- **REAL, DOUBLE**: Floating point number (approximate) with double-precision

# Numeric Data in SQL

- DECIMAL vs INT/FLOAT/DOUBLE

- FLOAT and DOUBLE are faster than DECIMAL
- DECIMAL values are exact
  - Example

floats: FLOAT	decimals: DECIMAL(3,2)
1.1	1.10
1.1	1.10
1.1	1.10

- `SELECT SUM(...)` → DECIMAL values are precise

SUM(floats)	SUM(decimals)
<u>3.3000000715255737</u>	<u>3.30</u>

# Temporal Data in SQL

- SQL Data Types

Time stamp?

**DATE**: 'YYYY-MM-DD'

- Range: 1000-01-01 to 9999-12-31
- E.g., '2020-03-01' for March 1, 2020

**TIME**: 'HH:MM:SS'

- Range: -838:59:59 to 838:59:59
- E.g., '14:30:03.5' for 3.5 seconds after 2:30pm

**DATETIME**: 'YYYY-MM-DD HH:MM:SS'

- Range: 1000-01-01 00:00:00 to 9999-12-31 23:59:59

**YEAR**: 'YYYY'

- Range: 1901 to 2155, or 0000 (illegal year values are converted to 0000)

Combination of DATE and TIME.

↑  
null value for Year.

# Temporal Data in SQL

- SQL Data Types

- **TIMESTAMP(*n*)**: **Unix time** (time since Jan 1, 1970)

- A way to track time as a **running total of seconds**
    - Range: **1970-01-01 00:00:01 UTC** to **2038-01-19 03:14:07 UTC**
    - Typically used for logging (keeping records of all the system events)
    - URL: <https://time.is/Unix>

Signed integer: 32 bits

Unixtime now:

**UNIX  
TIME  
SINCE 1970**

# Temporal Data in SQL

- SQL Data Types
  - **TIMESTAMP(*n*)**: Unix time (time since Jan 1, 1970)
    - Range: 1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
    - Typically used for logging (keeping records of all the system events)
  - Depending on size *n*, the display pattern changes

	Format
TIMESTAMP(14)	YYYYMMDDHHMMSS
TIMESTAMP(12)	YYMMDDHHMMSS
TIMESTAMP(10)	YYMMDDHHMM
TIMESTAMP(8)	YYYYMMDD
TIMESTAMP(6)	YYMMDD
TIMESTAMP(4)	YYMM
TIMESTAMP(2)	YY

# Large Objects in SQL

---

- SQL Data Types

- **BINARY(*n*)**: binary byte data type, with user-specified length *n*
  - Contains a byte strings, (rather than a character string)
  - Maximum length *n* = [0, 255]
- **VARBINARY(*n*)**: binary byte data type, with user-specified maximum length *n*
  - Maximum length *n* = [0, 65,535]
- **BLOB**: Binary Large OBject data type
  - TINYBLOB 0 – 255 bytes
  - BLOB 0 – 65,535 bytes (65 KB)
  - MEDIUMBLOB 0 – 16,777,215 bytes (16 MB)
  - LONGBLOB 0 – 4,294,967,295 bytes (4 GB)

# CREATE TABLE Construct

- A new relation is defined using the **CREATE TABLE** command:

**CREATE TABLE**  $r$

$(A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
 $(integrity-constraint_1),$   
 $\dots,$   
 $(integrity-constraint_k))$

- $r$  is the name of the relation,
- Each  $A_i$  is an attribute name in the schema of relation  $r$
- Each  $D_i$  is the data type of values in the domain of attribute  $A_i$

- Example:

**CREATE TABLE** instructor(  
    ID                            CHAR(5),  
    name                        VARCHAR(20),  
    dept\_name                    VARCHAR(20),  
    salary                      NUMERIC(8,2))

# Integrity Constraints in CREATE TABLE

- SQL prevents any update to the database that violates an **integrity constraint**
  - Integrity constraints allow us **to specify what data makes sense for us**
- Types of integrity constraints
  - Primary key: **PRIMARY KEY ( $A_1, \dots, A_n$ )**, *identifier of records in the table - if you have pk value.*
  - Foreign key: **FOREIGN KEY ( $A_m, \dots, A_n$ ) REFERENCES  $r$** , *you can specify the corresponding record.*
  - Unique key: **UNIQUE ( $A_1, \dots, A_n$ )**
  - Not null: **NOT NULL**
  - Value constraints: **CHECK (constraint)**, **DEFAULT**
- Example:

**CREATE TABLE** *instructor*(

<i>ID</i>	<b>CHAR(5),</b>
<i>name</i>	<b>VARCHAR(20) NOT NULL,</b>
<i>dept_name</i>	<b>VARCHAR(20),</b>
<i>salary</i>	<b>NUMERIC(8, 2),</b>
<b>PRIMARY KEY (<i>ID</i>),</b>	
<b>FOREIGN KEY (<i>dept_name</i>) REFERENCES <i>department</i>);</b>	

*in department table, there should be dept\_name attribute.*

# Declaring KEY and UNIQUE Constraints

- An attribute or list of attributes may be declared as PRIMARY KEY or UNIQUE.
  - Meaning: no two tuples of the relation may agree in all the attribute(s) on the list
    - That is, the attribute(s) do(es) **not allow duplicates** in values
    - PRIMARY KEY/UNIQUE can be used as an **identifier for each row**
  - Comparison: PRIMARY KEY vs UNIQUE.

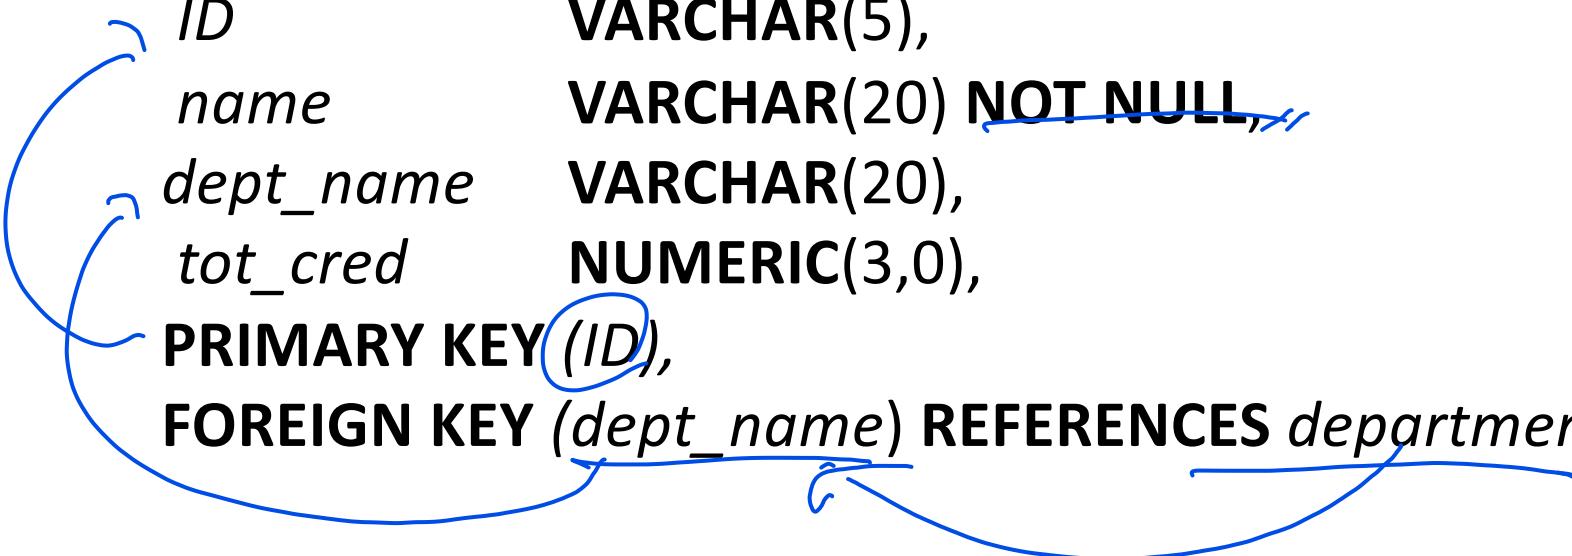
PRIMARY KEY	UNIQUE
Used to serve as a <u>unique identifier</u> for each row in a relation	<i>(can determine a record uniquely, but is still not a Pk)</i> <u>Uniquely determines a row</u> <b>which is not primary key</b> .
Cannot accept <u>NULL</u>	<u>Can accept NULL values</u> , (some DBMSs accept only one NULL value)
A relation can have <b>only one</b> primary key <i>PK just can be a combination of multiple attributes.</i> Clustered index	A relation can have <b>more than one</b> unique attributes Non-clustered index

# Examples

---

- **CREATE TABLE** *student* (

*ID*                    **VARCHAR(5),**  
*name*                **VARCHAR(20) ~~NOT NULL,~~**  
*dept\_name*        **VARCHAR(20),**  
*tot\_cred*           **NUMERIC(3,0),**  
**PRIMARY KEY** *(ID)*,  
**FOREIGN KEY** *(dept\_name)* **REFERENCES** *department*);



## Examples

---

- **CREATE TABLE** *student* (

*ID* **VARCHAR(5) PRIMARY KEY**,  
*name* **VARCHAR(20) NOT NULL**,  
*dept\_name* **VARCHAR(20)**,  
*tot\_cred* **NUMERIC(3,0)**,  
**FOREIGN KEY (dept\_name) REFERENCES department**);

↓ Simpler version.

# More Examples

- **CREATE TABLE** takes (

<i>ID</i>	<b>VARCHAR(5),</b>
<i>course_id</i>	<b>VARCHAR(8),</b>
<i>sec_id</i>	<b>VARCHAR(8),</b>
<i>semester</i>	<b>VARCHAR(6),</b>
<i>year</i>	<b>NUMERIC(4,0),</b>
<i>grade</i>	<b>VARCHAR(2),</b>

**PRIMARY KEY** (*ID, course\_id, sec\_id, semester, year*),

**FOREIGN KEY** (*ID*) **REFERENCES** student,

**FOREIGN KEY** (*course\_id, sec\_id, semester, year*)

**REFERENCES** section);

*a combination of attributes  
as a PK value.*

*foreign key can be  
also a combination of attributes*

## More Examples

---

- **CREATE TABLE** *course* (  
    *course\_id*      **VARCHAR**(8),  
    *title*            **VARCHAR**(50),  
    *dept\_name*      **VARCHAR**(20),  
    *credits*        **NUMERIC**(2,0),  
    **PRIMARY KEY** (*course\_id*),  
    **FOREIGN KEY** (*dept\_name*) **REFERENCES** *department*);

## More Examples

---

- **CREATE TABLE** course (  
    course\_id      **VARCHAR**(8),  
    title           **VARCHAR**(50),  
    dept\_name      **VARCHAR**(20) **DEFAULT** 'Comp. Sci',  
    credits        **NUMERIC**(2,0),  
    **PRIMARY KEY** (course\_id),  
    **FOREIGN KEY** (dept\_name) **REFERENCES** department);

If it is not set  
then follow: Default value.  
But what if user  
try to insert Null value?

# More Examples

---

- **CREATE TABLE** *neighbors*(  
    *name* **CHAR(30) PRIMARY KEY**,  
    *addr* **CHAR(50) DEFAULT '123 Sesame St.'**,  
    *phone* **CHAR(16)**);
- Inserting Elmo is a neighbor (inserted with the default value):
  - **INSERT INTO** *neighbors* (*name*)  
    **VALUES** ('Elmo');

<b>name</b>	<b>addr</b>	<b>phone</b>
'Elmo'	'123 Sesame St.'	NULL

# More Examples

---

- **CREATE TABLE** *neighbors*(  
    *name* **CHAR(30) PRIMARY KEY**,  
    *addr* **CHAR(50) DEFAULT '123 Sesame St.'**,  
    *phone* **CHAR(16) NOT NULL**);

- Inserting Elmo is a neighbor:
    - `INSERT INTO neighbors (name)  
VALUES ('Elmo');`
- If *phone* were NOT NULL, this insertion would have been **rejected**

# Integrity Constraints Recap

---

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
  - A single or multiple columns can be specified as a key
  - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID          NUMERIC(5,0),  
    name        VARCHAR(20),  
    city        VARCHAR(20),  
    state       CHAR(2),  
);
```

# Integrity Constraints Recap

---

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
  - A single or multiple columns can be specified as a key
  - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (
    ID          NUMERIC(5,0),
    name        VARCHAR(20),
    city        VARCHAR(20),
    state       CHAR(2),
    UNIQUE(name)
);
```



# Integrity Constraints Recap

---

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
  - A single or multiple columns can be specified as a key
  - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (
    ID          NUMERIC(5,0),
    name        VARCHAR(20),
    city        VARCHAR(20),
    state       CHAR(2),
    UNIQUE(name),
    UNIQUE(city, state),
);

```

*( ) can have multiple unique  
and declared multiple times.*

# Integrity Constraints Recap

---

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
  - A single or multiple columns can be specified as a key
  - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (
    ID          NUMERIC(5,0),
    name        VARCHAR(20),
    city        VARCHAR(20),
    state       CHAR(2),
    PRIMARY KEY(ID),
    UNIQUE(name),
    UNIQUE(city, state),
);
```

# Integrity Constraints Recap

---

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
  - A single or multiple columns can be specified as a key
  - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (  
    ID          NUMERIC(5,0) PRIMARY KEY,  
    name        VARCHAR(20) UNIQUE,  
    city        VARCHAR(20),  
    state        CHAR(2),  
    UNIQUE(city, state),  
);
```

# Integrity Constraints Recap

---

- Primary key, foreign key, and unique (candidate key) can be specified with DDL
  - A single or multiple columns can be specified as a key
  - Once a set of columns have been declared unique, any duplicate inputs are rejected

```
CREATE TABLE studio (
    ID          NUMERIC(5,0) PRIMARY KEY,
    name        VARCHAR(20) UNIQUE,
    city        VARCHAR(20),
    state       CHAR(2),
    UNIQUE(city, state),
    FOREIGN KEY (state) REFERENCES states (state)
    );

```

or states.state.

# Integrity Constraints Recap

---

- **NOT NULL** – disallowing null values
  - Null values indicate that the data is not known
  - These can cause problems in querying database
  - The Primary Key columns automatically prevent null being entered
  - *C.f.*, **NULL** – can be used to explicitly allow null values

```
CREATE TABLE studio (
    ID          NUMERIC(5,0) PRIMARY KEY,
    name        VARCHAR(20) NOT NULL,
    city        VARCHAR(20) NULL, It can explicitly say it can have null value
    state       CHAR(2) NOT NULL but not necessarily.
);
```

# Integrity Constraints Recap

- **DEFAULT** – A default value can be inserted in any column with this keyword

- E.g., **CREATE TABLE** movies(

movie\_title

release\_date

genre

)

**VARCHAR(40) NOT NULL,**

**DATE DEFAULT sysdate NULL,**

**VARCHAR(20) DEFAULT 'Comedy'**

**CHECK genre IN ('Comedy', 'Action', 'Drama')**

*SQL standard meaning  
Current time.*

*If you insert a new record with null value for release\_date  
then DBMS will automatically put sysdate for that.*

*then release\_date will never have null value? → No*

*→ You can update it to Null later.*

**VARCHAR(40) NOT NULL,**

**DATE DEFAULT CURRENT\_TIMESTAMP NULL,**

**VARCHAR(20) DEFAULT 'Comedy'**

**CHECK genre IN ('Comedy', 'Action', 'Drama')**

*MySQL doesn't have sysdate  
but this or Now()*

# Integrity Constraints Recap

- **CHECK** – Allows the inserted value to be checked

- *E.g., CREATE TABLE movies(*

<i>movie_title</i>	<b>VARCHAR(40) PRIMARY KEY,</b>
<i>release_date</i>	<b>DATE,</b>
<u><i>budget</i></u>	<b>INTEGER CHECK (<i>budget</i> &gt; 50000)</b>

)

- Table-level constraints can be defined; *E.g.,*

- **CREATE TABLE movies(**

<i>movie_title</i>	<b>VARCHAR(40) PRIMARY KEY,</b>
<i>release_date</i>	<b>DATE,</b>
<i>budget</i>	<b>INTEGER CHECK (<i>budget</i> &gt; 50000),</b>
<b>CONSTRAINT</b> <u><i>release_date_const</i></u>	
<b>CHECK</b> ( <i>release_date</i> <b>BETWEEN</b> '01-Jan-2000' <b>AND</b> '31-Dec-2009')	

)

↑  
check...

Constraint name. but why? → to be easy to maintain.  
easier to find. or compare.

# Table Updates (Updating Tuples)

- **INSERT**

- **INSERT INTO** *instructor*  
**VALUES** ('10211', 'Smith', 'Biology', 66000)

- **DELETE**

- **DELETE FROM** *student*

- Remove all tuples from the *student* relation
    - **TRUNCATE TABLE** *student*

*Remove all records* **But not a table structure.**



# Table Updates (Updating Table Schemas)

---

- **DROP**: Used to remove elements from a database, such as tables,
  - **DROP TABLE  $r$**  *both table structure and its all records are removed.*
  - Remove relation  $r$
  - C.f., **TRUNCATE (TABLE)  $r$**  is used to delete the data inside a table, but not the table itself.
- **ALTER**: Used to make changes to the table schema.
  - **ALTER TABLE  $r$  ADD  $A D$** 
    - $A$  is the name of the **new attribute** to add to relation  $r$ ;  $D$  is the **domain** of  $A$
    - All **existing tuples in the relation are assigned  $null$**  as the value for the new attribute
  - **ALTER TABLE  $r$  DROP  $A$** 
    - $A$  is the name of an **attribute** in  $r$
    - **Dropping of attributes** not supported by many databases (MySQL does)

# Table Updates (Updating Table Schemas)

- Examples

- **DROP TABLE** time\_slot\_backup;

↙ *Wearc entire table.*

- **ALTER TABLE** time\_slot\_backup **ADD** remark VARCHAR(20);

*Add new column.*

- **ALTER TABLE** time\_slot\_backup **MODIFY** remark CHAR(20);

*Type changed.*

- **ALTER TABLE** time\_slot\_backup **DROP** remark;

*delete a column.*



# Table Updates (Updating Table Schemas)

- Examples

- Drop a column
  - **ALTER TABLE** ~~salary~~ *instructor*  
**DROP COLUMN** *instructor*;

- **DROP** a PRIMARY KEY Constraint

- **ALTER TABLE** *instructor*  
**DROP PRIMARY KEY;**

it does not mean it drops the records of primary key  
but just take its primary key constraint off.

- **DROP** a FOREIGN KEY Constraint

- **ALTER TABLE** *instructor*  
**DROP FOREIGN KEY** *instructor\_ibfk\_1*;

- **DROP** DEFAULT

- **ALTER TABLE** *student*  
**ALTER** *tot\_cred* **DROP DEFAULT**;



# Table Updates (Updating Table Schemas)

---

- Examples
  - Drop a database
    - **DROP DATABASE** university;



# EOF

---

- Coming next:
  - Designing a database